



**BLAZE**  
INFORMATION SECURITY

**INTELLIGENT RISK PREVENTION**

<https://www.blazeinfosec.com>



June 4th 2018 • Julio Cesar Fort

**IT'S A TRAP!**

Remote detection of low & medium interaction honeypots

\$ whoami

Julio Fort

- Director of Professional Services at Blaze
- Been around the security community and industry for a while
- Lives in Kraków and quite likes it here

## Agenda



## **Presentation Roadmap**

**What will be discussed and what to expect**

# Agenda

- A brief overview of honeypots
- Prior art on the topic of detection
- Why are honeypots prone to detection?
- Case study: detecting a popular SSH honeypot
- Experiment: Internet-wide scanning for SSH honeypots
- Closing remarks

## On honeypots

---



**A brief overview of honeypots**

# Honeypots

- A honeypot *simulates*\* a real, vulnerable system
- The concept spans back a few decades, probably pioneered by Clifford Stoll (The Cuckoo's Egg)
- Considered an important asset for researchers doing threat intelligence
- Helps to learn more about tools, tactics and procedures of an adversary
- Gaining momentum lately with cyber deception technologies

\* actually, sometimes it is a real system

## Honeypots – different types

- Can be divided in three broad categories regarding their interaction levels:

### Interaction level



- Obviously, each comes with their own advantages and shortcomings
- The focus of this presentation is on the detection of **medium** interaction honeypots



# Honeypots

A green arrow pointing to the right, containing the word "Low" in white text.

Low

## Interaction level

- Provides basic functionality in an attempt to emulate a service
- Usually catches a specific exploit or a threat pattern (e.g., a certain worm or a particular type of attack)
- As a threat researcher you may probably not get much out of it
- Examples: honeyd, dionaea, spampot

# Honeypots

Medium

## Interaction level

- Enhanced interaction
- Pretty popular among the threat research community
- Easy to use, low maintenance and apparently it yields good results
- Examples: Kippo, Cowrie, Glustopf, Conpot

# Honeypots

High

## Interaction level

- Fully functioning system with extra instrumentation and logging capabilities
- **Advantages:** can be used to discover unknown vulnerabilities & exploits and are harder to detect
- **Disadvantages:** the compromised host can be used to attack others, high maintenance cost
- Examples: Sebek, custom VM + sensors

## PRIOR ART



About previous work on honeypot detection

## A word on prior art

- Work on this subject has been done in this past
- Academic work using SVM, etc. – not very practical
- Most stuff revolved around detecting virtual machines (useless these days)
- Sebek was detected and evaded on the fake Phrack 62 and 63 <sup>1</sup>
- Jon Oberhide & Manish Karir wrote a paper on how to detect honeyd <sup>2</sup>

<sup>1</sup> (<http://books.gigatux.nl/mirror/honeypot/final/ch09lev1sec2.html>)

<sup>2</sup> (<https://jon.oberheide.org/blog/2006/02/15/honeyd-remote-fingerprint/>)

## A word on prior art

- Tod Bearsley of Rapid7 and Andrew Morris did some work detecting Kippo from a pre-authenticated perspective - many of the techniques that were made public on detecting Kippo were patched in Cowrie 😞
- Most detection techniques for SSH honeypots rely on post-auth (within the simulated environment)
- In 2015 Darren Martyn discussed a couple of ways to detect Conpot in “OPSEC for honeypots”
- SHODAN introduced some of the fingerprinting techniques found by Martyn into the search engine.<sup>1</sup>

<sup>1</sup> <https://honeyscore.shodan.io>

## ABOUT DETECTION

---



**Why are many honeypots detection prone?**

## Honeypots – Proneness to detection

Low

**Interaction level**

Medium

**Interaction level**

- Nowhere close to fully emulate a real target system
- They have very little or incomplete features
- Or features may behave differently in a real system and in the emulated one



## Honey pots – Proneness to detection

High

### Interaction level

- Significantly harder to detect, as it is a real system
- Just realizing that you are inside a VM is so 2002 – no longer a reliable indicator of a honeypot (like discussed in the timeless presentation by GOBBLES’ “Wolves among us” at DEFCON 10)
- Still they may be prone to detection, like Sebek
- A live vulnerable system with additional instrumentation and logging can be a good idea, but maintenance is high.<sup>1</sup>

<sup>1</sup> <https://doublepulsar.com/eternalpot-lessons-from-building-a-global-nation-state-smb-exploit-honeypot-infrastructure-3f2a0b064ffe>

## GETTING PRACTICAL

---



**Case study: detecting Cowrie and Kippo**

## Case study: detecting Cowrie and Kippo

- Their emulated shell provides a not-very-complete bash terminal
- Many commands and common tools are missing too
- It's **pretty easy** to detect Kippo or Cowrie once you're inside the shell
- Also not difficult to detect them before falling into the trap

## Case study: detecting Cowrie and Kippo

Shell inside a real Linux system

```
julio@whatever2:~$ for((i=0; i<3; i++)); do echo "Is this real?"; done
Is this real?
Is this real?
Is this real?
julio@whatever2:~$ █
```

## Case study: detecting Cowrie and Kippo

Shell inside a Cowrie honeypot

```
root@svr04:~# for((i=0; i<3; i++)); do echo "Is this real?"; done
bash: for: command not found
bash: i: command not found
bash: i++: command not found
bash: do: command not found
bash: done: command not found
root@svr04:~#
```

Honeypot detection

# Detecting Kippo and Cowrie

How to detect them remotely without having to fall into the trap?

## Pre-authenticated detection strategy

- The basic idea is to understand the behavior of OpenSSH vs. the behavior of Kippo and Cowrie prior to authentication
- Pre-auth includes version advertisement, cipher suite and key exchange
- How does OpenSSH react to a given input – does Kippo/Cowrie handle it in the same way?
- Map the differences (if any) and take note of them
- Bottom-line: no need to waste time and resources brute forcing credentials for a fake system

## Some pre-authenticated detection techniques

- **Bad version:** advertising the client as supporting an invalid SSH version (e.g., SSH-31337.0) will make Kippo and Cowrie throw a “bad version” error
- **Double banner:** advertising the client with two banners of causes Cowrie to throw a “Packet corrupt” error.
- **Spacer:** advertising a client with 8 trailing “\n” characters causes Cowrie to throw a “Packet corrupt” error, and Kippo a “Protocol mismatch” response. OpenSSH, on the other hand, works just fine.
- **Fuzzing key exchange:** some corrupted payloads of key exchange, cipher suite and compression algorithms advertisement throw an exception on Kippo and Cowrie, but errors are handled gracefully in OpenSSH



# Pre-authenticated detection – Twisted

```
703 def _unsupportedVersionReceived(self, remoteVersion):
704     """
705     Called when an unsupported version of the ssh protocol is received from
706     the remote endpoint.
707
708     @param remoteVersion: remote ssh protocol version which is unsupported
709     by us.
710     @type remoteVersion: L{str}
711     """
712     self.sendDisconnect(DISCONNECT_PROTOCOL_VERSION_NOT_SUPPORTED,
713                        b'bad version ' + remoteVersion)
714
715
716 def dataReceived(self, data):
717     """
718     First, check for the version string (SSH-2.0-*). After that has been
719     received, this method adds data to the buffer, and pulls out any
720     packets.
721
722     @type data: L{bytes}
723     @param data: The data that was received.
724     """
725     self.buf = self.buf + data
726     if not self.gotVersion:
727         if self.buf.find(b'\n', self.buf.find(b'SSH-')) == -1:
728             return
729
730     # RFC 4253 section 4.2 ask for strict '\r\n' line ending.
731     # Here we are a bit more relaxed and accept implementations ending
732     # only in '\n'.
733     # https://tools.ietf.org/html/rfc4253#section-4.2
734     lines = self.buf.split(b'\n')
735     for p in lines:
736         if p.startswith(b'SSH-'):
737             self.gotVersion = True
738             # Since the line was split on '\n' and most of the time
739             # it uses '\r\n' we may get an extra '\r'.
740             self.otherVersionString = p.rstrip(b'\r')
741             remoteVersion = p.split(b'-')[1]
742             if remoteVersion not in self.supportedVersions:
743                 self._unsupportedVersionReceived(remoteVersion)
744             return
```

## Pre-authenticated detection – OpenSSH

```
451     chop(server_version_string);
452     debug("Local version string %.200s", server_version_string);
453
454     if (remote_major != 2 &&
455         !(remote_major == 1 && remote_minor == 99)) {
456         s = "Protocol major versions differ.\n";
457         (void) atomicio(vwrite, sock_out, s, strlen(s));
458         close(sock_in);
459         close(sock_out);
460         logit("Protocol major versions differ for %s port %d: "
461             "%.200s vs. %.200s",
462             ssh_remote_ipaddr(ssh), ssh_remote_port(ssh),
463             server_version_string, client_version_string);
464         cleanup_exit(255);
465     }
466 }
```

# Pre-authenticated detection – Twisted KEx

```
863     self.otherKexInitPayload = chr(MSG_KEXINIT) + packet
864     # This is useless to us:
865     # cookie = packet[: 16]
866     k = getNS(packet[16:], 10)
867     strings, rest = k[:-1], k[-1]
868     (kexAlgs, keyAlgs, encCS, encSC, macCS, macSC, compCS, compSC, langCS,
869      langSC) = [s.split(b',') for s in strings]
870     # These are the server directions
871     outs = [encSC, macSC, compSC]
872     ins = [encCS, macSC, compCS]
873     if self.isClient:
874         outs, ins = ins, outs # Switch directions
875     server = (self.supportedKeyExchanges, self.supportedPublicKeys,
876              self.supportedCiphers, self.supportedCiphers,
877              self.supportedMACs, self.supportedMACs,
878              self.supportedCompressions, self.supportedCompressions)
879     client = (kexAlgs, keyAlgs, outs[0], ins[0], outs[1], ins[1],
880              outs[2], ins[2])
881     if self.isClient:
882         server, client = client, server
883     self.kexAlg = ffs(client[0], server[0])
884     self.keyAlg = ffs(client[1], server[1])
885     self.nextEncryptions = SSHCiphers(
886         ffs(client[2], server[2]),
887         ffs(client[3], server[3]),
888         ffs(client[4], server[4]),
889         ffs(client[5], server[5]))
890     self.outgoingCompressionType = ffs(client[6], server[6])
891     self.incomingCompressionType = ffs(client[7], server[7])
892     if None in (self.kexAlg, self.keyAlg, self.outgoingCompressionType,
893                self.incomingCompressionType):
894         self.sendDisconnect(DISCONNECT_KEY_EXCHANGE_FAILED,
895                             b"couldn't match all kex parts")
896         return
897     if None in self.nextEncryptions.__dict__.values():
898         self.sendDisconnect(DISCONNECT_KEY_EXCHANGE_FAILED,
899                             b"couldn't match all kex parts")
900         return
```

## Pre-authenticated detection example

- Bad version example

```
root@research:~# echo -e "SSH-31337\r" | nc honeypot.whatever.io 22 | strings
SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
Protocol mismatch.
```

```
root@research:~# echo -e "SSH-31337\r" | nc honeypot.whatever.io 2022 | strings
SSH-2.0-OpenSSH_5.1p1 Debian-5
bad version 31337
root@research:~# echo -e "SSH-31337\r" | nc honeypot.whatever.io 2222 | strings
SSH-2.0-OpenSSH_6.0p1 Debian-4+deb7u2
bad version 31337
```

## PRE-AUTHENTICATED DETECTION

---



**LIVE DEMO**

## Pre-authenticated detection strategy

- Some false positives may arise: for example, sometimes incorrectly configured OpenSSH servers can throw “Packet corrupt” errors
- The script uses multiple techniques to detect a SSH honeypot and attributes score for each successful probe
- If score  $> 2$  it can be said with some degree of certainty it is a honeypot

## SCANNING



**Internet-wide scan to find SSH honeypots**

## INITIAL OBJECTIVES

---



- 
- ✓ **Map the entire Internet for Kippo and Cowrie honeypots**
  - ✓ **Find a suitable tool to allow us Internet-wide scanning**
  - ✓ **Plot them all on a map for visualization purposes**
-



## FINAL OBJECTIVES

---



**DETECT ALL THE THINGS!**



Tools of the trade

# Using masscan for our goals

A brief explanation on how to modify and use Robert Graham's masscan to serve our purposes

## Masscan

- Masscan has a not so well-known feature to actually send data to a server upon connecting to the port
- It also has a functionality for banner grabbing
- This combination suits our needs!

```
$ echo -e "SSH-HONEY\r\n" | base64  
U1NILUhPTkVZDQoK
```

```
./masscan --range 0.0.0.0/0 -p22 --banner -oL ssh_badversion.scan --exclude 255.255.255.255  
--rate 850000 --hello-string[22] U1NILUhPTkVZDQoK
```

## Modifying Masscan

modify proto-banner1.c from:

```
26: struct Patterns patterns[] = {  
27:     {"SSH-1.",          6, PROTO_SSH1, SMACK_ANCHOR_BEGIN},  
28:     {"SSH-2.",          6, PROTO_SSH2, SMACK_ANCHOR_BEGIN},
```

to:

```
struct Patterns patterns[] = {  
    {"SSH-H.",          6, PROTO_SSH1, SMACK_ANCHOR_BEGIN},  
    {"SSH-P.",          6, PROTO_SSH2, SMACK_ANCHOR_BEGIN},  
    {"HTTP/1.",         7, PROTO_HTTP, SMACK_ANCHOR_BEGIN},
```

# EXPERIMENT – INTERNET-WIDE SSH HONEYPOT SCANNING

---



## Notes about the experiment

- A cheap VPS will do the job pretty well



- It took ~20 hours to scan the entire Internet for honeypots



- Logs from banner had on average ~900mb



- Be prepared to handle dozens of abuse complaints per day



## Mass scanning methodology

- Masscan ports 22, 2022 and 2222 with the detection strings
- Triage the results using the response patterns we already learned
- Some results are obviously honeypots, others are only suspected (e.g., “Packet corrupt” errors may happen in other situations)
- Feed the suspected IPs into the detection script

# Results





## Unexpected results



## Unexpected results

- While scanning the Internet, a funny SSH banner appeared...

```
julio@whatever2:~/research/ssh-honey/old/mass-scans$ telnet 194.21.102.22
Trying 194.21.102.22...
Connected to 194.21.102.22.
Escape character is '^]'.
SSH-2.0-OpenSSH_5.8p2_hpn13v11 GCHQHONEYPOT34
^]
telnet> q
Connection closed.
```

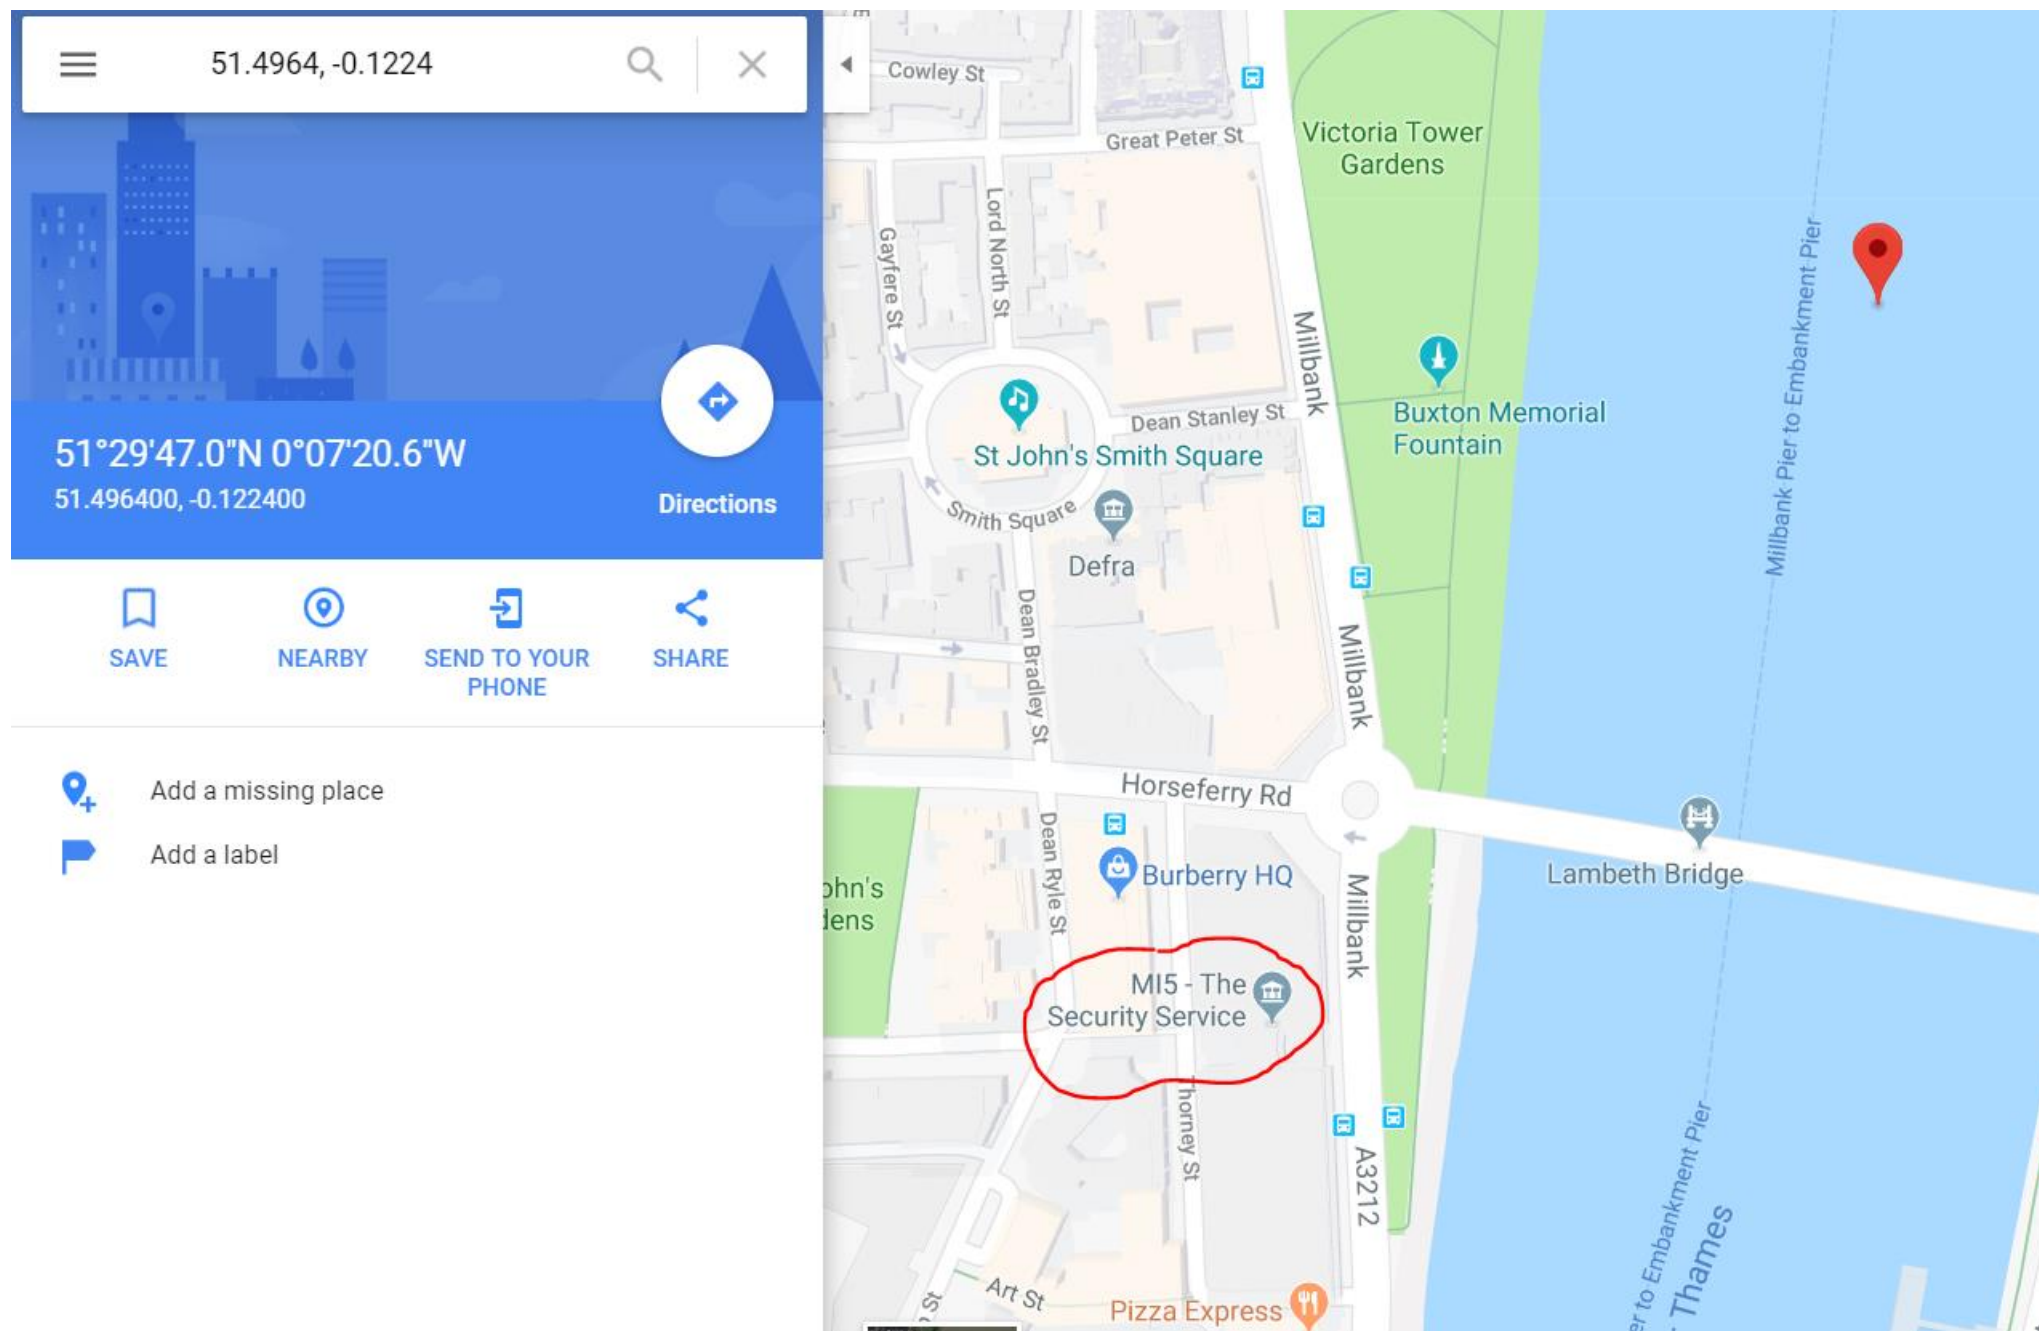
```
julio@whatever2:~/research/ssh-honey/old/mass-scans$ telnet 194.21.150.22
Trying 194.21.150.22...
Connected to 194.21.150.22.
Escape character is '^]'.
SSH-2.0-OpenSSH_5.2p1 gchqhoneypot12
```

## Unexpected results


### GeoIP2 City Results

IP Address	Country Code	Location	Postal Code	Approximate Coordinates*	Accuracy Radius	ISP	Organization
194.25.102	GB	United Kingdom, Europe		51.4964, -0.1224	200	High Availability Hosting Limited	High Availability Hosting Limited
194.25.150	GB	United Kingdom, Europe		51.4964, -0.1224	200	High Availability Hosting Limited	High Availability Hosting Limited







# Unexpected results




# Unexpected results




MI5 - The Security Service [GB] | <https://www.mi5.gov.uk/partnerships>





SECURITYSERVICE  
MI5

To report an imminent threat call **999** or ring the Anti-Terrorist Hotline on **0800 789 321**  
Current UK threat level: SEVERE [Read more about threat levels >>](#)

CAREERSWHO WE ARE ▾WHAT WE DO ▾HOW WE WORK ▾WHAT YOU CAN DO ▾NEWS

HOME / [HOW WE WORK](#) / [PARTNERSHIPS](#)

HOW WE WORK

Gathering intelligence +

Law and governance

Evidence and disclosure

Managing information +

Partnerships

National intelligence machinery

## PARTNERSHIPS

### Government departments

MI5 works particularly closely with the [Home Office](#) and has strong links with the [Foreign & Commonwealth Office](#), the [Cabinet Office](#), the [Northern Ireland Office](#), the [Department for Business, Innovation and Skills](#), and the [Ministry of Defence](#). We also advise all central government departments and agencies on protective security matters.

---

### Other intelligence agencies

We work very closely with both the [Secret Intelligence Service \(SIS\)](#) and the [Government Communications Headquarters \(GCHQ\)](#). We each have different but related functions, and there are many areas where we provide mutual assistance. We promote co-operation when a result can be achieved more effectively with the help of one or both of the other agencies. We are also a major customer for any intelligence they produce which is relevant to our functions and supplements our

# CLOSING REMARKS



## Closing remarks

- Honeypots are valuable tools for threat analysis and research...
- ... however, in many cases it is trivial to detect them
- Finding new detections & patching them will be an eternal cat-and-mouse game, with the attacker always winning
- Solely relying on honeypots for threat research is a flawed assumption
- A high interaction honeypot will certainly yield better results but it comes with a cost

Scripts and supporting material at our repository:

<https://github.com/blazeinfosec/detect-kippo-cowrie>





**Thank you!**

**Julio Cesar Fort • [julio@blazeinfosec.com](mailto:julio@blazeinfosec.com)**

**<https://www.blazeinfosec.com>**